

# mml87\_nb

*Kelvin*

*27 June 2017*

## MLE of NB parameters

The MLE estimations of NB parameters are as follows:

$$P(Y = y) = \frac{\text{count}(y)}{n}$$
$$P(X_j = x|Y = y) = \frac{\text{count}(x) \cup \text{count}(y)}{\text{count}(y)}$$

Notice that the log likelihood function that stated below is different from what mostly shown by others, which do not consider the case where  $Y_i = T$  and  $Y_i = F$ . Hence, it is not sure whether we can use the above MLE of parameters. The following is the implementations of MLE of NB parameters.

```
# y is the output node x is a set of input nodes estimated
# parameters are stored in a list smoothing is additive
# smoothing constant to avoid 0 probabilities but it can take
# any other values y parameters is always store at the end
mle_est_nb = function(y, x, data, smoothing = 1) {
  xIndices = which(colnames(data) %in% x)
  yIndex = which(colnames(data) == y)
  lst = list()
  for (i in 1:length(x)) {
    lst[[i]] = t((table(data[, c(yIndex, xIndices[i])]) +
      smoothing)/(rowSums(table(data[, c(yIndex, xIndices[i])])) +
      smoothing * nlevels(data[, xIndices[i]))))
  } # end for i
  # always store y parameters at the last
  lst[[length(x) + 1]] = (table(data[, yIndex]) + smoothing)/(nrow(data) +
    smoothing * nlevels(data[, yIndex]))
  names(lst) = c(x, y)
  return(lst)
}
```

## Negative log likelihood of NB

For a Naive Bayes model with binary variable, its parameters are  $\{P(Y_i = T), P(X_{ij}|Y_i = T), P(X_{ij}|Y_i = F)\}$ . To simply the notations, we use  $\{p_{i0}, p_{ij1}, p_{ij2}\}$  to denote the above probabilities respectively. The likelihood of Naive Bayes given a data set is then

$$l = \prod_i^n P(Y_i = T|\vec{X}_i)^{Y_i} (1 - P(Y_i = T|\vec{X}_i))^{1-Y_i}$$

where the posterior probability of  $Y_i$  given a vector  $\vec{X}_i = \langle X_{i1}, \dots, X_{im} \rangle$  is

$$\begin{aligned} P(Y_i = T | \vec{X}_i) &= \frac{P(Y_i = T) \prod_{j=1}^m P(X_{ij} | Y_i = T)}{P(\vec{X}_i)} \\ &= \frac{P(Y_i = T) \prod_{j=1}^m P(X_{ij} | Y_i = T)}{P(Y_i = T) \prod_{j=1}^m P(X_{ij} | Y_i = T) + (1 - P(Y_i = T)) \prod_{j=1}^m P(X_{ij} | Y_i = F)} \\ &= \frac{p_{i0} \prod_{j=1}^m p_{ij1}}{p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2}} \end{aligned}$$

The negative loglikelihood

$$\begin{aligned} L &= - \sum_{i=1}^n \left[ Y_i \ln p(Y_i | \vec{X}_i) + (1 - Y_i) \ln(1 - p(Y_i | \vec{X}_i)) \right] \\ &= - \sum_{i=1}^n \left[ Y_i \ln p_{i0} + (1 - Y_i) \ln(1 - p_{i0}) + Y_i \sum_{j=1}^m \ln p_{ij1} + (1 - Y_i) \sum_{j=1}^m \ln p_{ij2} - \ln \left( p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2} \right) \right] \end{aligned}$$

## NLL implementations

```
nll_auxiliary = function(dataPoint, pars, xIndices, yIndex) {
  ss = 0
  for (j in 1:(length(pars) - 1)) {
    ss = ss + log(p_ijk(dataPoint, pars, xIndices, xIndices[j],
      dataPoint[[yIndex]]))
  }
  ss = ss + log(pars[[length(pars)]] [dataPoint[[yIndex]]])
  return(ss)
}
```

## Fisher information matrix

The first derivatives of the above negative log likelihood w.r.t. each parameter are

$$\begin{aligned} \frac{\partial L}{\partial p_{i0}} &= - \sum_{i=1}^n \left[ \frac{Y_i}{p_{i0}} - \frac{1 - Y_i}{1 - p_{i0}} - \frac{\prod_{j=1}^m p_{ij1} - \prod_{j=1}^m p_{ij2}}{p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2}} \right] \\ \frac{\partial L}{\partial p_{ik1}} &= - \sum_{i=1}^n \left[ \frac{Y_i}{p_{ik1}} - \frac{p_{i0} \prod_{j=1}^m p_{ik1}}{p_{ik1} (p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})} \right] \\ \frac{\partial L}{\partial p_{ik2}} &= - \sum_{i=1}^n \left[ \frac{1 - Y_i}{p_{ik2}} - \frac{(1 - p_{i0}) \prod_{j=1}^m p_{ik2}}{p_{ik2} (p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})} \right] \end{aligned}$$

The second derivatives are

$$\begin{aligned}
\frac{\partial^2 L}{\partial p_{i0}^2} &= \sum_{i=1}^n \left[ \frac{Y_i}{p_{i0}^2} + \frac{1 - Y_i}{(1 - p_{i0})^2} - \left( \frac{\prod_{j=1}^m p_{ij1} - \prod_{j=1}^m p_{ij2}}{p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2}} \right)^2 \right] \\
\frac{\partial^2 L}{\partial p_{ik1}^2} &= \sum_{i=1}^n \left[ \frac{Y_i}{p_{ik1}^2} - \left( \frac{p_{i0} \prod_{j=1}^m p_{ij1}}{p_{ik1} (p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})} \right)^2 \right] \\
\frac{\partial^2 L}{\partial p_{ik2}^2} &= \sum_{i=1}^n \left[ \frac{1 - Y_i}{p_{ik2}^2} - \left( \frac{(1 - p_{i0}) \prod_{j=1}^m p_{ij2}}{p_{ik2} (p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})} \right)^2 \right] \\
\frac{\partial^2 L}{\partial p_{i0} \partial p_{ik1}} &= \sum_{i=1}^n \frac{\prod_{j=1}^m p_{ij1} p_{ij2}}{(p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})^2} \frac{1}{p_{ik1}} \\
\frac{\partial^2 L}{\partial p_{i0} \partial p_{ik2}} &= \sum_{i=1}^n \frac{\prod_{j=1}^m p_{ij1} p_{ij2}}{(p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})^2} \frac{-1}{p_{ik2}} \\
\frac{\partial^2 L}{\partial p_{ik1} \partial p_{ik2}} &= \sum_{i=1}^n \frac{\prod_{j=1}^m p_{ij1} p_{ij2}}{(p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2})^2} \frac{-p_{i0}(1 - p_{i0})}{p_{ik1} p_{ik2}}
\end{aligned}$$

Since FIM entries are expectations of the second derivatives, we need to take expectations for the first three second derivatives that contain  $Y_i$ . To simplify the notations, we use  $p_x$  to denote  $p_{i0} \prod_{j=1}^m p_{ij1} + (1 - p_{i0}) \prod_{j=1}^m p_{ij2}$ . Then the expectations become

$$\begin{aligned}
E \left( \frac{\partial^2 L}{\partial p_{i0}^2} \right) &= \sum_{i=1}^n \left[ \frac{\prod_{j=1}^m p_{ij1}}{p_{i0} p_x} + \frac{\prod_{j=1}^m p_{ij2}}{(1 - p_{i0}) p_x} - \left( \frac{\prod_{j=1}^m p_{ij1} - \prod_{j=1}^m p_{ij2}}{p_x} \right)^2 \right] \\
E \left( \frac{\partial^2 L}{\partial p_{ik1}^2} \right) &= \sum_{i=1}^n \left[ \frac{p_{i0} \prod_{j=1}^m p_{ij1}}{p_{ik1}^2 p_x} - \left( \frac{p_{i0} \prod_{j=1}^m p_{ij1}}{p_{ik1} p_x} \right)^2 \right] \\
E \left( \frac{\partial^2 L}{\partial p_{ik2}^2} \right) &= \sum_{i=1}^n \left[ \frac{(1 - p_{i0}) \prod_{j=1}^m p_{ij2}}{p_{ik2}^2 p_x} - \left( \frac{(1 - p_{i0}) \prod_{j=1}^m p_{ij2}}{p_{ik2} p_x} \right)^2 \right]
\end{aligned}$$

## FIM implementations

```

# p(x_ij/y=k) xIndices is a vector of x indices in data
# xIndex is the index of the particular x that we want
p_ijk = function(dataPoint, pars, xIndices, xIndex, yValue) {
  xValue = dataPoint[[xIndex]]
  xParsIndex = which(xIndices == xIndex)
  return(pars[[xParsIndex]][xValue, yValue])
}

# \prod_j p(x_ij/y=k) i is the dataPoint index j is the X
# index in data k is the value of Y, i.e. k \in [1,
# arity(y)]
prod_pijk = function(dataPoint, pars, xIndices, yValue) {
  mm = 1

```

```

for (i in 1:length(xIndices)) {
  xValue = dataPoint[[xIndices[i]]]
  mm = mm * pars[[i]][xValue, yValue]
}
return(mm)
}

# a function to calculate FIM
calculate_fim = function(prodPij1, prodPij2, px, probsMatrix,
  py1, py2, arities, yIndex) {
  # empty FIM
  fimDim = (arities[yIndex] - 1) + length(x) * arities[yIndex]
  fim = matrix(0, nrow = fimDim, ncol = fimDim)

  # off diagonal entries
  mm = prodPij1 * prodPij2/(px^2) # a common constant
  fim[1, -1] = colSums(mm/probsMatrix) # fill in 1st row of FIM
  fim[1, (2:ncol(fim))[odd(2:ncol(fim))]] = -1 * fim[1, (2:ncol(fim))[odd(2:ncol(fim))]]
  for (rowIndex in 2:(nrow(fim) - 1)) {
    if (ncol(fim) - rowIndex == 1) {
      fim[rowIndex, -(1:rowIndex)] = sum(-mm * py1 * py2/(probsMatrix[,
        rowIndex - 1] * probsMatrix[, -(1:(rowIndex -
        1))]))
    } else {
      fim[rowIndex, -(1:rowIndex)] = colSums(-mm * py1 *
        py2/(probsMatrix[, rowIndex - 1] * probsMatrix[,
        -(1:(rowIndex - 1))]))
    } # end else
  }
  fim = fim + t(fim) # duplicate upper to lower triangular fim

  # diagonal entries assume all variables are binary, hence the
  # diag[1] is always the 2nd derivative w.r.t. p_i0
  diag(fim)[1] = sum(prodPij1/(py1 * px) + prodPij2/(py2 *
    px) - ((prodPij1 - prodPij2)/px)^2)
  diag(fim)[-1] = colSums(py1 * py2 * prodPij1 * prodPij2/(px *
    probsMatrix)^2)

  return(fim)
}

```

## MML of Naive Bayes

$$I = -\ln K - \ln h(\vec{\theta}) + \frac{1}{2} \ln F(\vec{\theta}) - \ln f(D|\vec{\theta}) + \frac{|\vec{\theta}|}{2}$$

where  $\vec{\theta} = \langle p_{i0}, p_{ij1}, p_{ij2} \rangle, \forall j \in [1, m]$  is the set of parameters,  $|\vec{\theta}|$  is the number of free parameters,  $K$  is the lattice constant and  $h(\vec{\theta})$  is the parameter prior. A commonly used conjugate prior for binary variables is beta prior (i.e., beta distribution) with probability density function

$$f(x, \alpha, \beta) = \frac{x^{\alpha-1}(1-x)^{\beta-1}}{B(\alpha, \beta)}$$

where  $B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha+\beta)}$ . For simplicity, we assume all parameters are uniformly (i.e.,  $\alpha = \beta = 1$ ), hence for a single parameter prior is  $x(1-x)$ . Assuming all parameters are independent, we have

$$\ln h(\vec{\theta}) = \ln p_{i0} + \ln(1 - p_{i0}) + \sum_{j=1}^m [\ln p_{ij1} + \ln(1 - p_{ij1}) + \ln p_{ij2} + \ln(1 - p_{ij2})]$$

Substituting this into the above MML formula we get

$$I = - \left( \ln p_{i0} + \ln(1 - p_{i0}) + \sum_{j=1}^m [\ln p_{ij1} + \ln(1 - p_{ij1}) + \ln p_{ij2} + \ln(1 - p_{ij2})] \right) + \frac{1}{2} F(\vec{\theta}) - \ln f(D|\vec{\theta}) + \frac{d}{2} (1 + \ln k_d)$$

where  $d$  is the number of free parameters and  $k_d$  is the lattice constant for each free parameter.

## MML implementations

```
mml_NB = function(data, pars, fim, px, xIndices, yIndex, debug = FALSE) {
  # negative log likelihood
  nll = -sum(apply(data, 1, nll_auxiliary, pars = pars, xIndices = xIndices,
    yIndex = yIndex)) + sum(log(px))
  # log determinant of FIM
  logF = log(det(fim))
  # number of free parameters
  d = nrow(fim)
  # log prior
  logPrior = sum(unlist(lapply(pars, log)))
  # lattice constant
  k = c(0.083333, 0.080188, 0.077875, 0.07609, 0.07465, 0.07347,
    0.07248, 0.07163)
  if (d <= length(k)) {
    kd = k[d]
  } else {
    kd = min(k)
  }
  # mml
  l = -logPrior + 0.5 * logF + nll + 0.5 * d * (1 + log(kd))

  if (debug) {
    cat("mml=", l, "\n")
    cat("@ 1st part=", l - nll, "\n")
    cat("*** -logPrior=", -logPrior, "\n")
    cat("*** logFisher=", logF, "\n")
    cat("*** logLattice=", 0.5 * d * (1 + log(kd)), "\n")
    cat("@ 2nd part=nll=", nll, "\n")
  }

  return(l)
}
```

## Random model

```

dag = empty.graph(c("Y", paste0("X", 1:10)))
for (i in 2:nnodes(dag)) {
  dag = set.arc(dag, nodes(dag)[1], nodes(dag)[i])
}
cpts = randCPTs(dag, 2, 1)
n = 1000
data = rbn(cpts, n)
nVars = ncol(data)
arities = sapply(data, nlevels)
names(arities) = c()

```

## Executing MML\_NB

```

y = "Y"
x = c("X1", "X2", "X3", "X4", "X5", "X6")
yIndex = which(colnames(data) == y)
xIndices = which(colnames(data) %in% x)
# mle of parameters with smoothing
pars = mle_est_nb(y, x, data, 1)
# p(y=T)
py1 = pars[[length(pars)]] [[1]]
# p(y=F)
py2 = pars[[length(pars)]] [[2]]
# a vector of \prod_j p(x_ij/y=1)
prodPij1 = apply(data, 1, prod_pijk, pars = pars, xIndices = xIndices,
  yValue = 1)
# a vector of \prod_j p(x_ij/y=2)
prodPij2 = apply(data, 1, prod_pijk, pars = pars, xIndices = xIndices,
  yValue = 2)
# a vector of p_xi
px = py1 * prodPij1 + py2 * prodPij2
# a matrix of p(x_j/y=T) and p(y_j/y=F)
probsMatrix = c()
for (j in xIndices) {
  probsMatrix = cbind(probsMatrix, apply(data, 1, p_ijk, pars = pars,
    xIndices = xIndices, xIndex = j, yValue = 1), apply(data,
    1, p_ijk, pars = pars, xIndices = xIndices, xIndex = j,
    yValue = 2))
}
# FIM
fim = calculate_fim(prodPij1, prodPij2, px, probsMatrix, py1,
  py2, arities, yIndex)
# mml
(i_nb = mml_NB(data, pars, fim, px, xIndices, yIndex, debug = TRUE))

## mml= 340.7571
## @ 1st part= 59.18847
## *** -logPrior= 22.44771
## *** logFisher= 94.75266
## *** logLattice= -10.63557
## @ 2nd part=nll= 281.5687

```

```
## [1] 340.7571
```

## MML\_CPT

```
indexListPerNodePerValue = count_occurance(data, arities)
# forward_greedy_fast(data, indexListPerNodePerValue,
# arities, sampleSize, y, logFactorialSheet, base = exp(1),
# debug = T)
(i_cpt = mml_cpt(indexListPerNodePerValue, arities, sampleSize,
  xIndices, yIndex, logFactorialSheet, base = exp(1)))
```

```
## [1] 361.0366
```

The above mml + cpt is even smaller than mml + nb, which is expected to be shorter due to less number of parameters comparing with a full cpt!!!

## MML\_Logit

```
dataNumeric = factor2numeric(data)
vars = names(data)
sampleSize = n
(i_lr = mml_logit(data, arities, sampleSize, x, y, sigma = 3,
  debug = T))
```

```
## $mml
## [1] 296.2296
##
## $nlogPrior
## [1] 10.71211
##
## $nlogLattice
## [1] -5.685556
##
## $logF
## [1] 21.55309
##
## $nll
## [1] 280.4265
```

```
# forward_greedy(data, arities, vars, sampleSize, y, score =
# mml_logit, dataNumeric = dataNumeric, debug = T)
```

## Sanity check

This is a sanity check to ensure that nll is corrected calculated. The following code use gRain to compute the posterior probability  $P(Y_i|\vec{X}_i)$  based on the estimated cpts from data. The posterior probability given each data point is then used to compute the nll of the entire data set. The answer confirms that the above nll calculation is correct.

```
cptsEst = bn.fit(dag, data, method = "bayes")
obj = gRbase::compile(as.grain(cptsEst))
post = querygrain(obj, nodes = names(data)[c(yIndex, xIndices)],
```

```

    type = "conditional")
nll2 = 0
for (i in 1:nrow(data)) {
  x1 = data[i, "X1"]
  x2 = data[i, "X2"]
  x3 = data[i, "X3"]
  y = data[i, "Y"]
  nll2 = nll2 + log(post[x1, x2, x3, y])
}
-nll2

```

## Comparison

```
i_nb
```

```
## [1] 340.7571
```

```
i_cpt
```

```
## [1] 361.0366
```

```
i_lr
```

```
## $mml
```

```
## [1] 296.2296
```

```
##
```

```
## $nlogPrior
```

```
## [1] 10.71211
```

```
##
```

```
## $nlogLattice
```

```
## [1] -5.685556
```

```
##
```

```
## $logF
```

```
## [1] 21.55309
```

```
##
```

```
## $nll
```

```
## [1] 280.4265
```

## Observations

- The above tests show that nll for both naive bayes and logit are close to nll using true cpts. But mml\_nb is larger than mml\_logit and mml\_cpt even if the true model is a naive bayes model. This is likely to be caused by large message length for nb parameter priors and definitely large  $\log(\det(\text{fim}))$  as compared with mml\_logit. Could this be because of the mml\_logit is only for 1st order logit model, hence the model is simpler than a nb?
- In mml\_nb,  $\det(\text{fim})$  are almost always negative when X is a single variable. When Xs are two or more variables, it is less likely to have negative determinant. Not sure what's the problem. This problem may have been fixed. But occasionally the determinant is still a small negative number, perhaps due to under flow.